Configuring high availability

GitHub Enterprise Server supports a high availability mode of operation designed to minimize service disruption in the event of hardware failure or major network outage affecting the primary appliance.

About high availability configuration

In a high availability configuration, a fully redundant secondary GitHub Enterprise Server appliance is kept in sync with the primary appliance through replication of all major datastores.

Creating a high availability replica

In an active/passive configuration, the replica appliance is a redundant copy of the primary appliance. If the primary appliance fails, high availability mode allows the replica to act as the primary appliance, allowing minimal service disruption.

Monitoring a high-availability configuration

After configuration of high availability for your GitHub Enterprise Server instance, you can monitor the status of data replication among to your instance's replica nodes.

Initiating a failover to your replica appliance

You can failover to a GitHub Enterprise Server replica appliance using the command line for maintenance and testing, or if the primary appliance fails.

Recovering a high availability configuration

After failing over to a GitHub Enterprise Server appliance, you should regain redundancy as soon as possible rather than rely on a single appliance.

Removing a high availability replica

You can stop replication to a GitHub Enterprise Server replica temporarily, or permanently remove replication.

About geo-replication

Geo-replication on GitHub Enterprise Server uses multiple active replicas to fulfill requests from geographically distributed data centers.

